

CS57300  
PURDUE UNIVERSITY  
SEPTEMBER 27, 2021

---

# DATA MINING

# DECISION TREES

# TREE LEARNING

- ▶ Top-down recursive divide and conquer algorithm
  - ▶ Start with all training examples at root
  - ▶ Select **best** attribute/feature: Take a greedy view to decide how “good” an attribute is
  - ▶ Partition examples by selected attribute
  - ▶ Recurse and repeat
- ▶ Other issues:
  - ▶ When to stop growing
  - ▶ Pruning irrelevant parts of the tree

## CHOOSING AN ATTRIBUTE/FEATURE

- ▶ Be greedy: choose an attribute that can immediately minimize the misclassification rate (i.e., as if no further subtree will grow)
- ▶ A good feature splits the examples into subsets that distinguish among the class labels as much as possible... ideally into pure sets of "all positive" or "all negative"
- ▶ A good attribute leads to **highly certain** prediction for training examples sharing the same value on that attribute!

## MEASURING UNCERTAINTY: ENTROPY

- ▶ Used to quantify the amount of randomness of a probability distribution.
- ▶ Definition: Suppose a discrete random variable  $X$  has a distribution of  $P(X)$ . The **entropy**  $H(X)$  of  $X$  is defined by:

$$H(X) = - \sum_x p(x) \log_2 p(x)$$

## ENTROPY OF A RANDOM VARIABLE

A completely random binary variable with  $P(X)=[0.5,0.5]$  has entropy:

$$H(X) = -(0.5 \log_2 0.5 + 0.5 \log_2 0.5) = -(-0.5 + -0.5) = 1$$


A deterministic variable with  $P(X)=[1,0]$  has entropy:

$$H(X) = -(1 \log_2 1 + 0 \log_2 0) = -(0+0) = 0$$

A biased variable with  $P(X)=[0.75,0.25]$  has entropy:

$$H(X) = 0.8113$$

The entropy of a probability distribution **p** expresses the **amount of uncertainty** that we have about the values of  $X$



## MEASURING CHANGE OF UNCERTAINTY: INFORMATION GAIN

- ▶ How much does a feature split decrease the entropy?

$$Gain(S, A) = \underline{Entropy(S)} - \sum_{v \in values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no


Entropy(S)

$$= -9/14 \log 9/14 - 5/14 \log 5/14$$

$$= 0.9400$$

# INFORMATION GAIN

$$Gain(S, A) = Entropy(S) - \sum_{v \in values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$



$$\begin{aligned} Entropy(\text{Income}=\text{High}) \\ = -2/4 \log 2/4 - 2/4 \log 2/4 = 1 \end{aligned}$$

$$\begin{aligned} Entropy(\text{Income}=\text{Med}) \\ = -4/6 \log 4/6 - 2/6 \log 2/6 = 0.9183 \end{aligned}$$


$$\begin{aligned} Entropy(\text{Income}=\text{Low}) \\ = -3/4 \log 3/4 - 1/4 \log 1/4 = 0.8113 \end{aligned}$$

$$\begin{aligned} Gain(S, \text{Income}) \\ = 0.9400 - (4/14 [1] + 6/14 [0.9183] + 4/14 [0.8113]) \\ = 0.029 \end{aligned}$$




## EXAMPLE: CHOOSE THE ATTRIBUTE WITH LARGEST INFORMATION GAIN

- ▶ Features: color, shape, size
- ▶ What's the best feature to use at root?



## EXAMPLE: CHOOSE THE ATTRIBUTE WITH LARGEST INFORMATION GAIN

Example	Color	Shape	Size	Class
1	Red	Square	Big	+
2	Blue	Square	Big	+
3	Red	Circle	Big	+
4	Red	Circle	Small	-
5	Green	Square	Small	-
6	Green	Square	Big	-



- ▶  $H(S) = -(0.5 \cdot \log 0.5 + 0.5 \cdot \log 0.5) = 1$
- ▶  $\text{Gain}(S, \text{Color})$   
 $= 1 - 0.5 \cdot (-0.67 \cdot \log 0.67 - 0.33 \cdot \log 0.33) - 0.17 \cdot (-1 \cdot \log 1 - 0 \cdot \log 0) - 0.33 \cdot (-1 \cdot \log 1 - 0 \cdot \log 0) = 0.54$
- ▶  $\text{Gain}(S, \text{Shape}) = 1 - 0.67 \cdot (-0.5 \cdot \log 0.5 - 0.5 \cdot \log 0.5) - 0.33 \cdot (-0.5 \cdot \log 0.5 - 0.5 \cdot \log 0.5) = 0$
- ▶  $\text{Gain}(S, \text{Size}) = 1 - 0.67 \cdot (-0.75 \cdot \log 0.75 - 0.25 \cdot \log 0.25) - 0.33 \cdot (-1 \cdot \log 1 - 0 \cdot \log 0) = 0.46$

## BUILDING TREE RECURSIVELY

**Buildtree**(examples, attributes)

*/\*examples: a list of training examples at the current node  
attributes: a set of candidate attributes to place question on\*/*

**If** examples={} **then** return

**If** examples have the same label  $y$  **then** return a leaf node with label  $y$

**If** attributes={} **then** return a leaf node with the majority label in examples

$A = \text{Best\_attribute}(\text{examples}, \text{attributes})$  */\*Suppose attribute  $A$  has  $n$  possible values\*/*

Create an internal node,  $\text{node}(A)$ , with  $n$  children

**For** attribute  $A$ 's  $i$ -th possible value  $A(i)$ :

The  $i$ -th child of  $\text{node}(A) = \text{Buildtree}(\{\text{examples with its value on } A \text{ being } A(i)\}, \text{attributes}-\{A\})$

## DEALING WITH CONTINUOUS ATTRIBUTES

- ▶ Discretize the value of a continuous attribute into several intervals
  - ▶ Two bins for the continuous variable  $x$ :  $x \leq t$ ,  $x > t$
  - ▶  $\text{Gain}(S, x, t) = \text{Entropy}(S) - |S_{x \leq t}| * \text{Entropy}(S_{x \leq t}) / |S| - |S_{x > t}| * \text{Entropy}(S_{x > t}) / |S|$
  - ▶  $\text{Gain}(S, x) = \max_t \text{Gain}(S, x, t)$

## ADDITIONAL ATTRIBUTE SELECTION CRITERIA: GINI GAIN


- ▶ Similar to information gain
- ▶ Uses gini index instead of entropy

$$Gini(X) = 1 - \sum_x p(x)^2$$

- ▶ Measures decrease in gini index after split:

$$Gain(S, A) = Gini(S) - \sum_{v \in values(A)} \frac{|S_A|}{|S|} Gini(S_A)$$

# COMPARING INFORMATION GAIN TO GINI GAIN

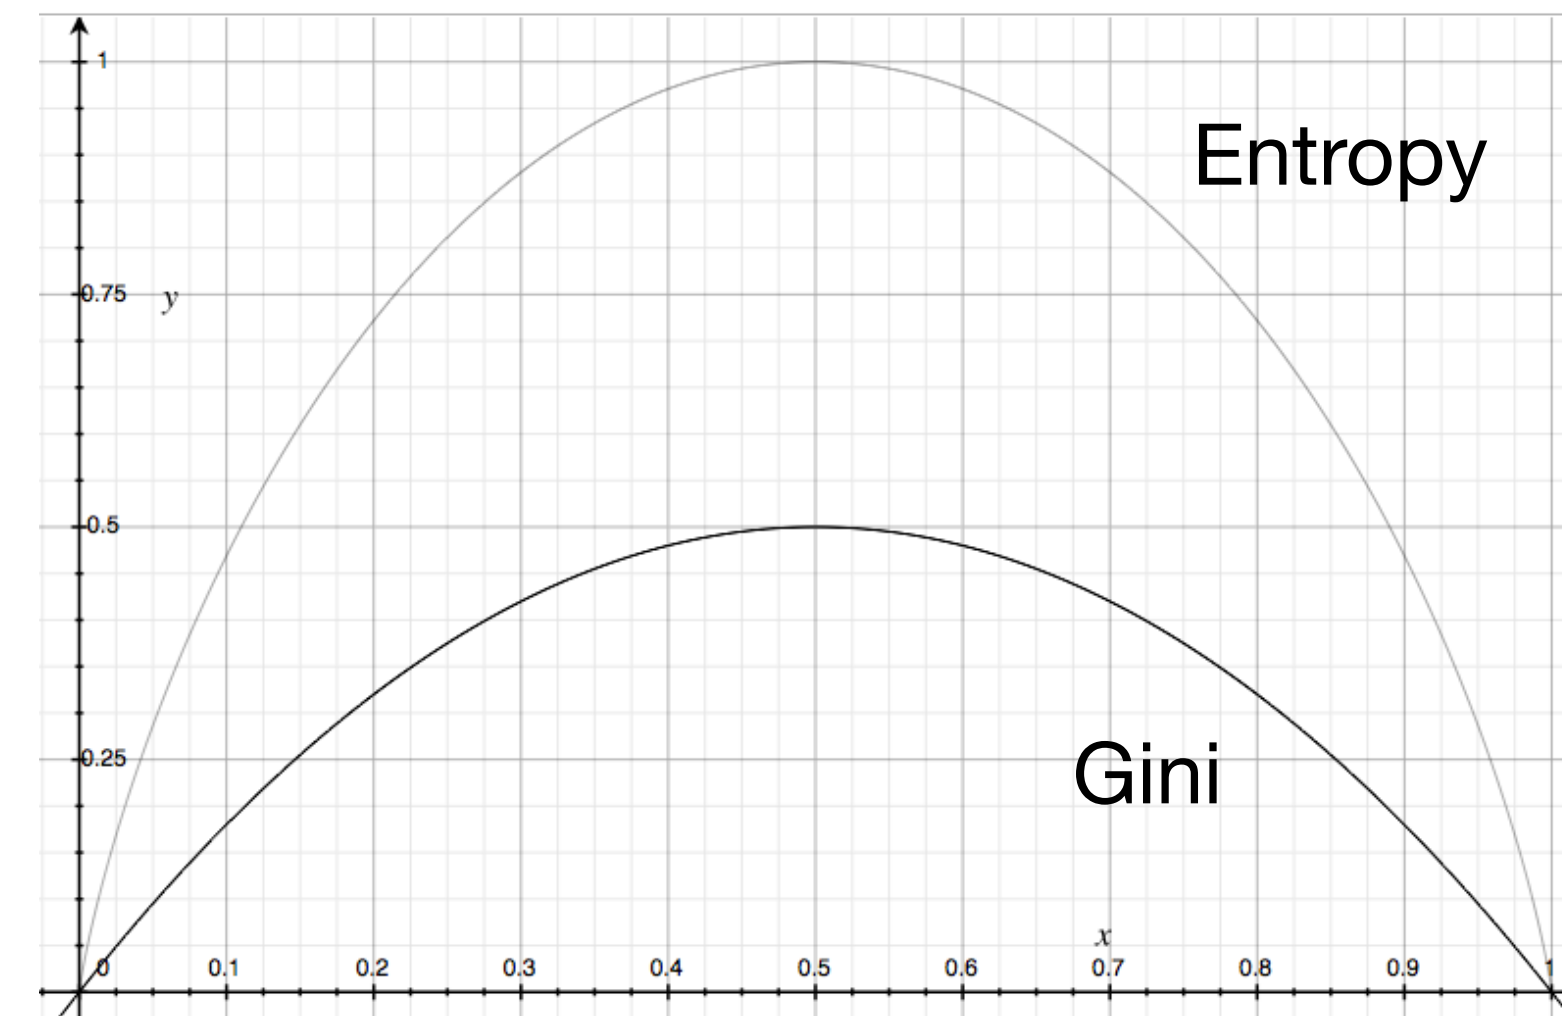



$$IG = 0$$

$$GG = 0$$



# COMPARING INFORMATION GAIN TO GINI GAIN



$$IG = 1.0 - \left[ \frac{2}{12} 0 \right] - \left[ \frac{4}{12} 0 \right] - \left[ \frac{6}{12} 0.919 \right] = 0.541$$

$$GG = 0.5 - \left[ \frac{2}{12} 0 \right] - \left[ \frac{4}{12} 0 \right] - \left[ \frac{6}{12} 0.444 \right] = 0.278$$

## ADDITIONAL ATTRIBUTE SELECTION CRITERIA: CHI-SQUARE SCORE

- ▶ Widely used to test independence between two categorical attributes (e.g., feature and class label)
- ▶ Considers counts in a contingency table and calculates the normalized squared deviation of observed (actual) values from expected (predicted) values

$$\chi^2 = \sum_{i=1}^k \frac{(o_i - e_i)^2}{e_i}$$

- ▶ Sampling distribution is known to be chi-square distributed, given that cell counts are above minimum thresholds



# CALCULATING EXPECTED VALUES FOR A CELL

$$\chi^2 = \sum_{i=1}^k \frac{(o_i - e_i)^2}{e_i}$$

$$o_{(0,+)} = a$$

$$e_{(0,+)} = p(A = 0, C = +) \cdot N$$

$$= p(A = 0)p(C = +|A = 0) \cdot N$$

$$= p(A = 0)p(C = +) \cdot N$$

$$= \left[ \frac{a + b}{N} \right] \cdot \left[ \frac{a + c}{N} \right] \cdot N$$

Class

	+	-
0	a	b
1	c	d

**N**

Attribute

(assuming independence)

## EXAMPLE CALCULATION

		<b>Observed</b>				<b>Expected</b>	
		Buy	No buy			Buy	No buy
High		2	2	High		2.57	1.43
Med		4	2	Med		3.86	2.14
Low		3	1	Low		2.57	1.43

$$\begin{aligned}\chi^2 &= \sum_{i=1}^k \frac{(o_i - e_i)^2}{e_i} = \left( \frac{(2 - 2.57)^2}{2.57} \right) + \left( \frac{(4 - 3.86)^2}{3.86} \right) + \left( \frac{(3 - 2.57)^2}{2.57} \right) \\ &\quad + \left( \frac{(2 - 1.43)^2}{1.43} \right) + \left( \frac{(2 - 2.14)^2}{2.14} \right) + \left( \frac{(1 - 1.43)^2}{1.43} \right) \\ &= 0.57\end{aligned}$$

## WHEN TO STOP GROWING

- ▶ Full growth methods
  - ▶ There are no examples left
  - ▶ All examples at a node belong to the same class
  - ▶ There are no attributes left for further splits
- ▶ What impact does this have on the quality of the learned trees?
  - ▶ Trees **overfit** the training data and accuracy on testing data suffers

## OVERFITTING


- ▶ Consider a distribution  $D$  of data representing a population and a sample  $D_S$  drawn from  $D$ , which is used as training data
- ▶ Given a model space  $M$ , a score function  $S$ , and a learning algorithm that returns a model  $m \in M$ , the algorithm **overfits** the training data  $D_S$  if:  
 $\exists m' \in M$  such that  $S(m, D_S) > S(m', D_S)$  but  $S(m, D) < S(m', D)$
- ▶ In other words, there is another model ( $m'$ ) that is better on the entire distribution and if we had learned from the full data we would have selected it instead

## HOW TO AVOID OVERFITTING IN DECISION TREES

- ▶ Post-pruning
  - ▶ Separate the training data into a training set and a validation set (i.e., a pruning set).
  - ▶ Fully grow a tree
  - ▶ Use the pruning set to evaluate the utility of pruning (i.e. deleting) nodes from the tree
- ▶ Pre-pruning
  - ▶ Apply a statistical test to decide whether to expand a node
  - ▶ Add penalty terms in scoring functions to prefer trees with smaller sizes


## POST-PRUNING METHOD: REDUCED ERROR PRUNING

- ▶ Grow a full tree  $T$  using the **training set**
- ▶ Let  $v$  be an internal node of the current tree  $T$
- ▶ If we prune at  $v$  to create a new tree  $T'$ , in  $T'$ , the subtree rooted at  $v$  will be replaced by a leaf node  $v'$ , whose label is the majority label for all **training examples** fall under  $v$
- ▶ Define the gain of pruning at  $v$  as, in the **pruning set**, # of misclassified examples under  $v$  (in  $T$ ) - # of misclassified examples that in  $v'$  (in  $T'$ )
- ▶ Repeat: Prune at node with largest gain until only negative gain nodes remain



## REDUCED ERROR PRUNING EXAMPLE

### Training set



Gain at node  $X_2$  (left):

$$3-2=1$$


Gain at node  $X_2$  (right):

$$1-3=-2$$

Gain at node  $X_1$ :  $4-5=-1$


Prune at node  $X_2$  (left)!

### Pruning set



## REDUCED ERROR PRUNING EXAMPLE

## Training set



Gain at node  $X_2$ :  $1-3=-2$

Gain at node  $X_1$ :  $3-5=-2$

Done!

## Pruning set

